

Finding Missed Compiler Optimizations by Differential Testing

Gergő Barany
Inria Paris, France
`gergo.barany@inria.fr`

CC 2018

Main takeaways

Does your compiler *always* optimize well?

- ▶ compare compilers' outputs to find missed optimizations
- ▶ automated toolchain finds minimal test cases
- ▶ issues found in GCC, Clang, CompCert:
 - ▶ peephole optimizations, dead stores, useless spills, missed instruction selection patterns, missed copy coalescing, ...

Example: missing range analysis

Generated source code:

```
int f(int p, int q) {  
    return q + (p % 6) / 9;  
}
```

($p \% 6 \in [-5, 5]$,
division truncates to 0)

Clang:

```
movw r2, #43691  
movt r2, #10922  
smmul r2, r0, r2  
add r2, r2, r2, lsr #31  
add r2, r2, r2, lsl #1  
sub r0, r0, r2, lsl #1  
movw r2, #36409  
movt r2, #14563  
smmul r0, r0, r2  
asr r2, r0, #1  
add r0, r2, r0, lsr #31  
add r0, r0, r1
```

GCC:

```
mov r0, r1
```

https://bugs.llvm.org/show_bug.cgi?id=34517 (fixed)

Example: redundant code

Source code:

```
int fn3(  
    double c,  
    int *p, int *q)  
{  
    int i = (int)c;  
    *p = i;  
    *q = i;  
    return i;  
}
```

Clang:

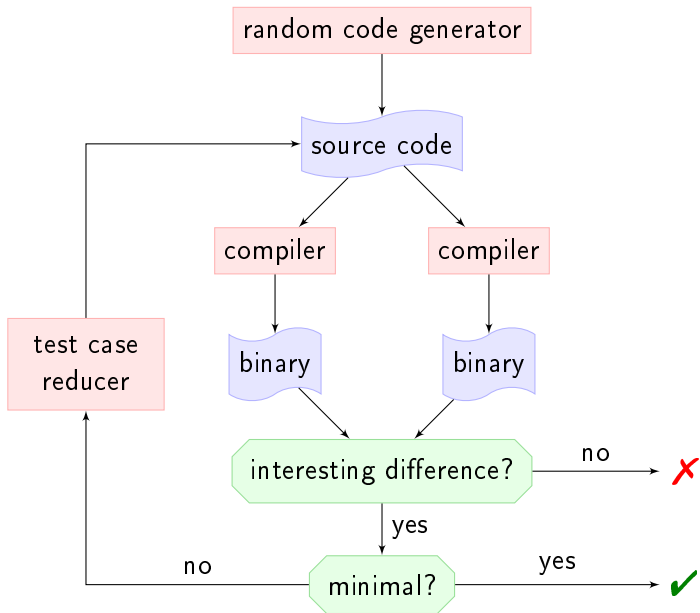
```
vcvt.s32.f64 s2, d0  
vstr s2, [r0]  
vcvt.s32.f64 s2, d0  
vcvt.s32.f64 s0, d0  
vmov r0, s0  
vstr s2, [r1]
```

GCC:

```
vcvt.s32.f64 s15, d0  
vstr.32 s15, [r0]  
vmov r0, s15  
vstr.32 s15, [r1]
```

https://bugs.llvm.org/show_bug.cgi?id=33199 (fixed)

Randomized differential testing



Randomized differential testing for missed optimizations

random code generator

off-the-shelf tools: Csmith, ldrngen (or many others)

test case reducer

off-the-shelf tool: C-Reduce

interesting difference?

custom tool: optdiff

- ▶ binary analysis to find optimization differences
- ▶ assigns scores to binaries, compares

optdiff

- ▶ based on angr binary analysis framework
 - ▶ multi-platform (x86, x86-64, ARM, PowerPC, ...)
 - ▶ Python API
- ▶ load binary, compute CFG, estimate basic block frequencies w_b

Checkers: local scoring functions $c : \textit{instruction} \rightarrow \mathbb{N}$

Total score:
$$s = \sum_{b \in f} w_b \cdot \sum_{i \in b} c(i)$$

Examples: number of instructions, general memory loads/stores, stack loads/stores, function calls, floating-point arithmetic instructions, vector instructions, ...

Checker implementation: instructions

Checkers

- ▶ Python functions with `@checker` decorator
- ▶ inspect one instruction at a time

@checker

```
def instructions(arch, instr):  
    """Number of instructions."""  
    return 1
```


Checker implementation: loads

@checker

```
def loads(arch, instr):  
    """Number of memory loads."""  
    op = instr.insn.mnemonic  
    if is_arm(arch):  
        if op == 'ldrd': # load doubleword  
            return 2  
        elif re.match('ldm.*', op): # load multiple  
            return len(instr.insn.operands)-1  
        return bool(re.match('v?ldr.*', op)) # load one word  
    ... # other architectures
```

Example: useless spill

Source code:

```
char fn2(  
    float p)  
{  
    char c=(char)p;  
    return c;  
}
```

Clang:

```
vcvt.u32.f32 s0, s0  
vmov r0, s0
```

instruction score: 2
stack load score: 0

GCC:

```
vcvt.u32.f32 s15, s0  
sub sp, sp, #8  
vstr.32 s15, [sp, #4]  
ldrb r0, [sp, #4]  
add sp, sp, #8
```

instruction score: 5
stack load score: 1

https://gcc.gnu.org/bugzilla/show_bug.cgi?id=80861
(confirmed, diagnosed)

CompCert: an example

Source code:

```
int fn10(int p1) {  
  int a, b, c, d, e, v, f;  
  a = 0;  
  b = c = 0;  
  d = e = p1;  
  v = 4;  
  f = e * d | a * p1 + b;  
  return f;  
}
```

CompCert:

```
str r4, [sp, #8]  
mov r4, #0  
mov r12, #0  
mov r1, r0  
mov r2, r1  
mul r3, r2, r1  
mla r2, r4, r0, r12  
orr r0, r3, r2  
ldr r4, [sp, #8]
```

- ▶ dead code `v = 4;` causes spilling
- ▶ missed copy coalescing
- ▶ missed constant propagation and folding: `a * p1 + b = 0`

Undefined behavior: the good

Undefined behavior may be compiled arbitrarily

— do we have to be careful?

Unproblematic cases:

- ▶ Clang and GCC treat many cases identically, comparisons OK
- ▶ `char f(float p) { return (char) p; }`
(assume never called with bad values of `p`)
- ▶ `x < x + 1` \rightarrow true
(undefined for `x` signed integer)

Undefined behavior: the bad

Problematic cases:

- ▶ **unconditional** undefined behavior, e. g.,
`int fn(int a) { int x = 0; return a / x; }`
- ▶ **infinite loops**:
`while (x) { y = ...; }`
- ▶ C-Reduce likes to produce such cases
- ▶ no compiler warnings but different 'optimized' code

Workarounds

- ▶ static analysis to find UB/nontermination? **ineffective** 😞
- ▶ accept some cases
- ▶ don't let random generators produce loops/problematic constructs

Why *randomized* differential testing?

Arguments against random input programs

- ▶ examples look artificial
- ▶ may not correspond to real-world performance problems

Advantages of random input programs

- ▶ unlimited amount of code available
- ▶ controlled sublanguage (loop-free, only constructs of interest)
- ▶ reducer output looks artificial even for real-world input

Results

Some missed optimizations found

	reported	fixed	other
GCC	6	1	13
Clang	3	3	3
CompCert	3	3	6

- ▶ generally treated as low priority by developers
- ▶ many duplicates

Causes

- ▶ missing/wrong rules or costs
- ▶ phase ordering
- ▶ weak heuristics

Summary

- ▶ compare compilers' outputs to find missed optimizations
- ▶ automated toolchain finds minimal test cases
- ▶ issues found in GCC, Clang, CompCert

<https://github.com/gergo-/missed-optimizations>

Thank you for your attention

This research was partially supported by ITEA 3 project no. 14014, ASSUME.