

Finding Missed Compiler Optimizations

Gergő BARANY - Inria Paris, France



...By differential testing

even tiny examples show optimization differences:

Source	Clang	GCC
<pre>int f(int p, int q) { return q + (p % 6) / 9; }</pre>	<pre>movw r2, #43691 movt r2, #10922 smmul r2, r0, r2 add r2, r2, r2, lsr #31 add r2, r2, r2, lsl #1 sub r0, r0, r2, lsl #1 movw r2, #36409 movt r2, #14563 smmul r0, r0, r2 asr r2, r0, #1 add r0, r2, r0, lsr #31 add r0, r0, r1 bx lr</pre>	<pre>mov r0, r1 bx lr</pre>

Clang did not know that $(p \% 6) / 9 = 0$ on ints, fixed now

Method

- * generate random C programs
- * compile with different compilers
- * **compare**: count instructions of interest
- * reduce to minimal example

...or more directly

Source	GCC
<pre>int g(short p, double q) { int a = 0; if (p) a = (int) q; return a; }</pre>	<pre>cmp r0, #0 sub sp, sp, #8 < useless stack frame vcvtne.s32.f64 s15, d0 streq r0, [sp, #4] < useless spill vmovne r0, s15 vstrne.32 s15, [sp, #4] < useless spill add sp, sp, #8 < useless pop bx lr</pre>

dead stores never reloaded, don't correspond to anything in the source

found By **liveness analysis** on the Binary

some results

tested GCC, Clang, and CompCert

found missed arithmetic optimizations in each

and cases of: Bad spilling, Bad coalescing, dead stores, redundant computations, missing instruction selection patterns

some reported, some fixed

more examples, preprint

<https://github.com/gergo-/missed-optimizations>